

MovieLens Project

Julien Geffroy

11.02.2022

A. Create train and validation sets

```
#####  
# Create edx set, validation set (final hold-out test set)  
#####  
  
# Note: this process could take a couple of minutes  
  
if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")  
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")  
if(!require(data.table)) install.packages("data.table", repos = "http://cran.us.r-project.org")  
  
library(tidyverse)  
library(caret)  
library(data.table)  
  
# MovieLens 10M dataset:  
# https://grouplens.org/datasets/movielens/10m/  
# http://files.grouplens.org/datasets/movielens/ml-10m.zip  
  
dl <- tempfile()  
download.file("http://files.grouplens.org/datasets/movielens/ml-10m.zip", dl)  
  
ratings <- fread(text = gsub(":", "\t", readLines(unzip(dl, "ml-10M100K/ratings.dat"))),  
  col.names = c("userId", "movieId", "rating", "timestamp"))  
  
movies <- str_split_fixed(readLines(unzip(dl, "ml-10M100K/movies.dat")), "\\:", 3)  
colnames(movies) <- c("movieId", "title", "genres")  
  
# if using R 3.6 or earlier:  
movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(levels(movieId))[movieId],  
  title = as.character(title),  
  genres = as.character(genres))  
  
# if using R 4.0 or later:  
movies <- as.data.frame(movies) %>% mutate(movieId = as.numeric(movieId),  
  title = as.character(title),  
  genres = as.character(genres))
```

```

movielens <- left_join(ratings, movies, by = "movieId")

# Validation set will be 10% of MovieLens data
set.seed(1, sample.kind="Rounding") # if using R 3.5 or earlier, use `set.seed(1)`
test_index <- createDataPartition(y = movielens$rating, times = 1, p = 0.1, list = FALSE)
edx <- movielens[-test_index,]
temp <- movielens[test_index,]

# Make sure userId and movieId in validation set are also in edx set
validation <- temp %>%
  semi_join(edx, by = "movieId") %>%
  semi_join(edx, by = "userId")

# Add rows removed from validation set back into edx set
removed <- anti_join(temp, validation)
edx <- rbind(edx, removed)

rm(dl, ratings, movies, test_index, temp, movielens, removed)

```

B. Data Exploration:

⇒ see Quiz: MovieLens Dataset

C. Data transformation

#Edx dataset transformation:userId and movieId should be treat as factors for some analysis purposes.

```

edx.copy <- edx
edx.copy$userId <- as.factor(edx.copy$userId)
edx.copy$movieId <- as.factor(edx.copy$movieId)

```

#SparseMatrix function is used in order to get an output Of sparse matrix of class dgCMatrix.

To use this function, the userId & movieId are converted to numeric vectors.

```

edx.copy$userId <- as.numeric(edx.copy$userId)
edx.copy$movieId <- as.numeric(edx.copy$movieId)
sparse_ratings <- sparseMatrix(i = edx.copy$userId,
j = edx.copy$movieId ,
x = edx.copy$rating,
dims = c(length(unique(edx.copy$userId)),
length(unique(edx.copy$movieId))),
dimnames = list(paste("u", 1:length(unique(edx.copy$userId)), sep = ""),
paste("m", 1:length(unique(edx.copy$movieId)), sep = "")))

```

#Remove the copy created

```
rm(edx.copy)
```

#The first 10 users

```
sparse_ratings[1:10,1:10]
```

```
## 10 x 10 sparse Matrix of class "dgCMatrix"
## [[ suppressing 10 column names 'm1', 'm2', 'm3' ... ]]
##
## u1 .....
## u2 .....
## u3 .....
## u4 .....
## u5 1.....3...
## u6 .....
## u7 .....
## u8 .2.5..34....
## u9 .....
## u10 .....3...
```


```
##Convert rating matrix into a recommenderlab sparse matrix
rate_Mat <- new("realRatingMatrix", data = sparse_ratings)
rate_Mat
```

69878 x 10677 rating matrix of class 'realRatingMatrix' with 9000055 ratings.

D. Similarity measures

```
#calculate the user similarity using the cosine similarity
similarity_users <- similarity(rate_Mat[1:50,],
method = "cosine",
which = "users")
```

```
image(as.matrix(similarity_users), main = "similarity of Users")
```




#Using the same approach, compute similarity between movies.

```
similarity_movies <- similarity(rate_Mat[,1:50],  
method = "cosine",  
which = "items")
```

```
image(as.matrix(similarity_movies), main = "similarity of Movies")
```

similarity of Movies



E. Dimension reduction

```
#implicitly restarted Lanczos bidiagonalization algorithm (IRLBA)  
set.seed(1, sample.kind = "Rounding")
```

```
## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding' sampler  
## used
```

```
Y <- irlba(sparse_ratings, tol=1e-4, verbose=TRUE, nv = 100, maxit = 1000)
```


```
## Working dimension size 107
```

```
## Initializing starting vector v with samples from standard normal distribution.  
## Use 'set.seed' first for reproducibility.
```

```
## irlba: using fast C implementation
```

```
# plot singular values  
plot(Y$d, pch=20, col = "blue", cex = 1.5, xlab='Singular Value', ylab='Magnitude',  
main = "User-Movie Matrix")
```

User-Movie Matrix



```
# calculate sum of squares of all singular values  
all_sing_val <- sum(Y$d^2)
```

```
# variability described by first 6, 12, and 20 singular values  
first_six <- sum(Y$d[1:6]^2)  
print(first_six/all_sing_val)
```

```
## [1] 0.6187623
```

```
first_twl <- sum(Y$d[1:12]^2)  
print(first_twl/all_sing_val)
```

```
## [1] 0.7052297
```


```
first_twt <- sum(Y$d[1:20]^2)  
print(first_twt/all_sing_val)
```

```
## [1] 0.7646435
```

```
perc_vec <- NULL  
for (i in 1:length(Y$d)) {  
  perc_vec[i] <- sum(Y$d[1:i]^2) / all_sing_val  
}
```

```
plot(perc_vec, pch=20, col = "blue", cex = 1.5, xlab='Singular Value', ylab='% Sum of Squares of Singular
lines(x = c(0,100), y = c(.90, .90))
```

k for Dimensionality Reduction



```
#value of K
k = length(perc_vec[perc_vec <= .90])
k
```

```
## [1] 55
```

```
#get the decomposition of Y ; matrices U, D, and V
U_k <- Y$u[, 1:k]
dim(U_k)
```

```
## [1] 69878 55
```

```
D_k <- Diagonal(x = Y$d[1:k])
dim(D_k)
```

```
## [1] 55 55
```

```
V_k <- t(Y$v)[1:k, ]
dim(V_k)
```

```
## [1] 55 10677
```

F. Relevant Data

#1. Determine the minimum number of movies per user.

```
min_no_movies <- quantile(rowCounts(rate_Mat), 0.9)
print(min_no_movies)
```

```
## 90%
```

```
## 301
```

#2. Determine the minimum number of users per movie.

```
min_no_users <- quantile(colCounts(rate_Mat), 0.9)
print(min_no_users)
```

```
## 90%
```

```
## 2150.2
```

#3. Select the users and movies matching these criteria.

```
rate_movies <- rate_Mat[rowCounts(rate_Mat) > min_no_movies,
colCounts(rate_Mat) > min_no_users]
rate_movies
```

```
## 6978 x 1068 rating matrix of class 'realRatingMatrix' with 2313148 ratings.
```

Movie effect

#calculate the average of all ratings of the edx dataset

```
mu <- mean(edx$rating)
```

#calculate b_i on the training dataset

```
movie_m <- edx %>%
group_by(movieId) %>%
summarize(b_i = mean(rating - mu))
```

predicted ratings

```
predicted_ratings_bi <- mu + validation %>%
left_join(movie_m, by='movieId') %>%
.$b_i
```

Movie and user effect

#calculate b_u using the training set

```
user_m <- edx %>%
left_join(movie_m, by='movieId') %>%
group_by(userId) %>%
summarize(b_u = mean(rating - mu - b_i))
```

#predicted ratings


```

predicted_ratings_bu <- validation %>%
left_join(movie_m, by='movieId') %>%
left_join(user_m, by='userId') %>%
mutate(pred = mu + b_i + b_u) %>%
.$pred

```

Movie, user and time effect

```

#create a copy of validation set , valid, and create the date feature which is the timestamp converted to
valid <- validation
valid <- valid %>%
mutate(date = round_date(as_datetime(timestamp), unit = "week"))

```

```

#calculate time effects ( b_t) using the training set
temp_m <- edx %>%
left_join(movie_m, by='movieId') %>%
left_join(user_m, by='userId') %>%
mutate(date = round_date(as_datetime(timestamp), unit = "week")) %>%
group_by(date) %>%
summarize(b_t = mean(rating - mu - b_i - b_u))

```

```

#predicted ratings
predicted_ratings_bt <- valid %>%
left_join(movie_m, by='movieId') %>%
left_join(user_m, by='userId') %>%
left_join(temp_m, by='date') %>%
mutate(pred = mu + b_i + b_u + b_t) %>%
.$pred

```

The root mean square error (RMSE) models for movies, users and time effects

```

#calculate the RMSE for movies
rmse_model_1 <- RMSE(validation$rating,predicted_ratings_bi)
rmse_model_1

```

```
## [1] 0.9439087
```

```

#calculate the RMSE for users
rmse_model_2 <- RMSE(validation$rating,predicted_ratings_bu)
rmse_model_2

```

```
## [1] 0.8653488
```

```

#calculate the RMSE for time effects
rmse_model_3 <- RMSE(valid$rating,predicted_ratings_bt)
rmse_model_3

```

```
## [1] 0.8652511
```


From the movie and user effects combined, our RMSE decreased by almost 10% with respect to the only movie effect. The improvement on the time effect is not significant, (about a decrease of 0.011%). The regularization would be performed using only the movie and user effects.

```
#remove valid before regularization  
rm(valid)
```

G. Regularization

```
## remembering that lambda is a tuning parameter. We can use cross-validation to choose it  
lambdas <- seq(0, 10, 0.25)  
rmsees <- sapply(lambdas, function(l){  
  mu_reg <- mean(edx$rating)  
  b_i_reg <- edx %>%  
    group_by(movieId) %>%  
    summarize(b_i_reg = sum(rating - mu_reg)/(n()+1))  
  b_u_reg <- edx %>%  
    left_join(b_i_reg, by="movieId") %>%  
    group_by(userId) %>%  
    summarize(b_u_reg = sum(rating - b_i_reg - mu_reg)/(n()+1))  
  predicted_ratings_b_i_u <-  
    validation %>%  
    left_join(b_i_reg, by = "movieId") %>%  
    left_join(b_u_reg, by = "userId") %>%  
    mutate(pred = mu_reg + b_i_reg + b_u_reg) %>%  
    .$pred  
  return(RMSE(validation$rating,predicted_ratings_b_i_u))  
})
```

```
qplot(lambdas, rmsees)
```



The optimal lambda for the full model

#For the full model, the optimal $\hat{\lambda}$ is given as

```
lambda <- lambdas[which.min(rmses)]
lambda
```

```
## [1] 5.25
```

```
rmse_model_4 <- min(rmses)
rmse_model_4
```

```
## [1] 0.864817
```

Summary of the rmse on validation set for Linear regression models

#summarize all the rmse on validation set for Linear regression models

```
rmse_results <- data.frame(methods=c("movie effect", "movie + user effects", "movie + user + time effects",
```

```
kable(rmse_results) %>%
kable_styling(bootstrap_options = "striped", full_width = F, position = "center") %>%
kable_styling(bootstrap_options = "bordered", full_width = F, position = "center") %>%
column_spec(1, bold = T) %>%
column_spec(2, bold = T, color = "white", background = "#D7261E")
```

methods	rmse
movie effect	0.9439087
movie + user effects	0.8653488
movie + user + time effects	0.8652511
Regularized Movie + User Effect Model	0.8648170

The regularization gets down the RMSE's value to 0.8648170.

POPULAR , UBCF and IBCF algorithms of the recommenderlab package

#POPULAR algorithms of the recommenderlab package

```
model_pop <- Recommender(rate_movies, method = "POPULAR",
param=list(normalize = "center"))
```

#prediction example on the first 10 users

```
pred_pop <- predict(model_pop, rate_movies[1:10], type="ratings")
as(pred_pop, "matrix")[1:10]
```

Rmse for popularity based recommender engine

```
set.seed(1, sample.kind = "Rounding")
```

```
## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding' sampler
## used
```

#Calculation of rmse for popular method

```
eval <- evaluationScheme(rate_movies, method="split", train=0.7, given=-5)
```

#ratings of 30% of users are excluded for testing

```
model_pop <- Recommender(getData(eval, "train"), "POPULAR")
```

```
prediction_pop <- predict(model_pop, getData(eval, "known"), type="ratings")
```

```
rmse_pop <- calcPredictionAccuracy(prediction_pop, getData(eval, "unknown"))[1]
rmse_pop
```

```
## RMSE
```

```
## 0.8482917
```

User based cosine factorization recommender engine

#Estimating rmse for UBCF using Cosine similarity and selected n as 50 based on cross-validation

```
set.seed(1, sample.kind = "Rounding")
```

```
## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding' sampler
```

```
## used
```

```
model <- Recommender(getData(eval, "train"), method = "UBCF",  
param=list(normalize = "center", method="Cosine", nn=50))  
  
prediction <- predict(model, getData(eval, "known"), type="ratings")  
  
rmse_ubcf <- calcPredictionAccuracy(prediction, getData(eval, "unknown"))[1]  
rmse_ubcf
```

```
## RMSE  
## 0.8589153
```

Item based Cosine factorization recommender engine

```
#Estimating rmse for IBCF using Cosine similarity and selected n as 350 based on cross-validation  
set.seed(1, sample.kind = "Rounding")
```

```
## Warning in set.seed(1, sample.kind = "Rounding"): non-uniform 'Rounding' sampler  
## used
```

```
model <- Recommender(getData(eval, "train"), method = "IBCF",  
param=list(normalize = "center", method="Cosine", k=350))  
prediction <- predict(model, getData(eval, "known"), type="ratings")  
rmse_ibcf <- calcPredictionAccuracy(prediction, getData(eval, "unknown"))[1]  
rmse_ibcf
```


```
## RMSE  
## 0.963769
```

Rmse from popularity, user and item based recommender engine

```
rmse_crossvalidation <- data.frame(methods=c("Popularity-Based model", "User-based Model", "Item-  
Based Model"), kable(rmse_crossvalidation) %>%  
kable_styling(bootstrap_options = "striped", full_width = F, position = "center") %>%  
kable_styling(bootstrap_options = "bordered", full_width = F, position = "center") %>%  
column_spec(1, bold = T) %>%  
column_spec(2, bold = T, color = "white", background = "#D7261E")
```

methods	rmse
Popularity-Based model	0.8482917
User-based Model	0.8589153
Item-Based Model	0.9637690

```
plot(rmse_crossvalidation, annotate = 1, legend = "topleft")
title("ROC curve, Three Models")
```



H. Regularization

The best model recommended from this research from the validation set is the regularized Movie + User effect Model with the least root mean square of 0.8648170.